

C++

Basics

- `for(auto v: vArray)` - range over an iterable without explicit index
- `enum class` - Doesn't explicitly map to int, safer and recommended
- `func(&int)` - Reference to int, like pointer but dereferencing happens automatically and introduces some restrictions similar to how the `const` keyword restricts modification, recommended instead of pointers
- If one needs to reassign, use a pointer. References can only be set once.
- `int* myPointer = &someInt; std::cout << *myPointer;`. The order of the `const` keyword(s) is very important to determine if the target or the pointer itself is `const`. This doesn't prevent the target from being changed some other way, only not through that pointer
- Always pair unions with an `enum` to represent which type its supposed to take on
- Properties of an object accessed with `object.property`, of a pointer to an object using `pointer->property` (equivalent to `(*pointer).property`), or simply `.` if it is a reference & not raw pointer, and members of a class are accessed using `std::cout`
- `lvalue` - an object that occupies some identifiable location in memory and can be assigned to
- `rvalue` - expressions that aren't `lvalues`, an object that isn't in memory, usually on the RHS in expression
- Expression evaluation may be short circuited (compiler optimization). For example: `1 != 2 && 6/3 == 1` will skip the right hand side division, unless it contains an assignment or something that could affect the control flow.
- The iteration expression is executed after the condition and the loop body have been executed, unlike JS
- The Turing Halting problem is the reason why not all programs (even those without external inputs) cannot be checked for errors / termination without executing the entire program (hence run-time errors exist)
- `switch` - case statements fall through unless `break` is used.
- Value types such as `int`, `double`, `bool` are passed by values to functions - new copies of them are created on the stack and the original values cannot be accessed directly from within the function scope
- Pre / postconditions specify the valid domain / expected range of a function (for example when dividing by an argument, it must not equal 0). These can be simply written as a comment `// PRE: condition` or expressed using `assert (e>=0 || b!=0)` by importing `#include <cassert>`, a sufficiently significant and easy to express pre / post condition is worth while expressing programatically.
- Functions are only visible after they have been declared, which can be problematic if two functions call each other. A solution is to firstly declare a function `int g();` defining its return and argument types, from which point it can be used (it's in scope), despite the implementation `int g() {}` being somewhere later in the file.
- Pointers are not initialized with `nullptr` by default to expose bugs. If this is desired they should be manually assigned in the empty constructor

Containers

- Vectors are resizable array wrappers and usually used since they're far more flexible. Similar to `List<Type>` in C#
- `vector.at(index)` allows safe index access, throwing an error if the given index is out of bounds. A compiler flag can be enabled so regular indexing has the same effect instead of simply reading the memory location.

- Most standard library functions treat indexes as uints, which can lead to underflow problems when used in operations with regular ints
- Long type names can be aliased: `using matrix = std::vector<std::vector<int>>;`, although `auto` may be more appropriate in one-off cases.
- Vectors should be passed by reference `func(std::vector<int>& numbers)` to avoid copying lots of data to the stack. Arrays are automatically passed by reference, but I'm not sure about their exact functionality yet
- Const references are especially useful for vectors, preventing the copying of the data but avoiding modification of the original array
- In modern C++, a for each loop can be used `for(int number: numbers)`
- Container is the overarching interface, each child class of which has the same iterator implementation. Another example is the `std::unordered_set<T>`, which offers the same behavior as a mathematical set
- When working with const containers, the `container::const_iterator` must be used instead. Of course, the const iterator doesn't allow modifying elements
- **Linked lists** don't require blocks of area and elements can be stored anywhere on the heap. They are linked through pointers to one another, allowing efficient modification of elements in the middle of the list but slower indexing (each element's pointer must be traversed starting from the top). `std::list<T>` is the standard implementation

Modularity

- Rather than including source `.cpp` files, using header files ensures that the "library" is only compiled once. Alternatively it allows calling functions from pre-compiled `.obj` files (for example from a closed-source library)
- All source files referenced are compiled separately, with non-main files being compiled into an `.obj` file with their implementations. A linker then copies implementations into the missing usages to create the final binary.
- Use the same header file for implementation and usage
- Errors are thrown to allow the user of a library to decide how to handle unexpected cases
- Namespaces group functions to prevent duplicate naming issues. They are referenced as follows: `namespace_name::func_name()`
- An entire namespace can be imported (functions available without prefix) through using `namespace` but this obscures the origin of a function and is recommend against.

OOP

- Use `struct` for plain data structures without any access modifiers or OOP features, otherwise use richer objects of a `class`. Struct properties are public by default, class props are private but can be controlled
- Struct assignment `myType item = a;` copies the members of `a`, unlike JS
- complex operator `+(complex a, complex b) { return a+=b; }` overloading default operators. This also includes implementation for type casting (for example `double` from a rational number class)
- The `const` modifier after a functions argument indicates functions that const initializations of the object may call. All other functions are inaccessible / private to const objects.
- `this` is usually a const reference to the current instance, its properties can be accessed directly by their name as shorthand for `this->property`
- The constructor can be called directly `myClass obj (arg1, arg2);` or simply `myClass obj;` (which calls the empty argument constructor, which is automatically generated to initialize properties with their null values if no such constructor has been written)

- Property initialization directly from the constructor arguments can be written using the following shorthand: `rational(int n, int d): num(n), den(d)` rather than writing `this.n = n; ...`
- `new` - Assigns memory on the heap for the object and returns a pointer. Has to be explicitly deleted (even after it leaves scope). Useful to allow a stack variable to be accessed by its pointer from outside of the scope in which it was created (otherwise it'll be automatically deleted).
- Classes incorporating dynamic memory may need to define a unique destructor `~ClassName()` which handles the deletion of dynamic properties to prevent memory leaks after the object itself has left scope or `delete` is called upon it. The default generated destructor simply calls the destructor of all properties, but this often doesn't suffice for complex data structures like a linked list
- representation - the properties / variables of a class, what stores memory
- abstract class, similar to an interface in Go, simply a collection of methods such a class must implement, can be used to specify what an argument is expected to have. Implemented as `class Implementor: public Abstract {}`, this is **inheritance**
- Polymorphism - one interface used to represent many other types which may satisfy it
- `virtual` - May be redefined later in a derived class, `virtual void x = 0` means it **must** be redefined otherwise the class cannot be instantiated, there is no default implementation.
- Base functions / properties can be accessed within subclass implementations
- Calling `delete` on an abstract object calls the destructor of the shallowest subclass (as it has access to the most "additional" properties)
- `dynamic_cast` can be used to check what derived class an abstract argument is
- Resource handle - A class that is responsible for managing underlying resources, these provide custom copy implementations to prevent violating validity, for example assigning a vector to another variable results in two vectors that refer to the **same** underlying elements. Such handles should implement a **copy constructor** `ClassName(const ClassName& a)` so underlying resources are correctly initialized in a different place in memory (instead of two handles pointing to the same underlying data unintentionally in case the only explicit property is a pointer to heap data)
- Assignment to an already existing object can be handled with **copy assignment**, simply overriding the default assignment operator `Shape& operator=(const Shape&)`
- One should never explicitly call the destructor, this becomes invalid
- Marking a constructor `explicit` prevents automatic type conversion
- Marking a default function (destructor, copy constructor / assignment) with `=delete` creates a compile time error if they're unintentionally used

Smart Pointers

- These allow a stack variable to "own" the object it points to, automatically deleting it once the smart pointer leaves the scope and will no longer be used
- Nearly always recommended instead of explicitly calling `new` and `delete`, these are theoretically never needed
- It is essential to use heap memory sometimes, for example underlying structs in a pointer vector. In such a case, using smart pointers is advantageous, potentially avoiding the need for custom destructors / copy operators.
- Although using stack only (memory freed once leaves scope) is often simpler, sometimes a factory-type architecture is advantageous for creating objects, in which case returning a smart pointer is your best bet for transfer of ownership.
- Use the `auto smartP = make_unique<MyClass>(constructor args)` or `make_shared` functions, which calls `new` under the hood

- Further pointers pointing the same object can be created simply by calling the copy constructor: `shared_pointer<myClass> sharedP = otherSharedP`, the underlying object is only deleted once all shared pointers leave scope
- Regular dereferencing operators `*` and `->method()` can be called on smart pointers

Dynamic Arrays

- A so-called heap array is allocated by calling `new Type[num of elements]`, which allocated a block of **contiguous** (as opposed to a linked list without random access) memory large enough to contain all those elements and returns a pointer to the first element. The `std::vector` type uses dynamic arrays.
- **Pointer Arithmetic** manipulates pointers with respect to their type. For example `int* intPointer = 0x...; cout << (intPointer + 1) - intPointer` would print 4 (or whatever the size in bytes of pointers on that machine). Array indexing does this under the hood
- `delete[] firstElementPointer;` liberates the entire allocated memory of the dynamic array

Generics

- Prefixing a class / function with `template<typename T>` accepts a type as a generic argument, so that `T` can be used throughout implementations
- So-called function objects can be defined by implementing the `()` operator, for example `bool operator()(const T& x) const { return x<val; }`
- Type aliases can be defined: `using value_type = T` is a public property of all container classes in the standard library, accessed using `Class::value_type`

Floating-point numbers

Floating point number systems are how types such as `double` and `float` represent real number approximations. Such systems allow storing and working with numbers in vastly different orders of magnitude and are denoted as follows:

$$F^*(\beta, p, e_{\min}, e_{\max})$$

$$\pm \sum_{i=0}^{p-1} (d_i \beta^{-i}) \times \beta^e$$

$$d_0.d_1, \dots, d_{p-1} \times \beta^e$$

Where the digits are called the *mantissa* and the exponent e indicates the order of magnitude, as in scientific notation 1.6×10^{-19} :

- β - base (for example 10 - decimal or 2 - binary)
- p - precision, how many significant figures are used to represent the mantissa
- $e_{\min/\max}$ - the range of possible exponents / orders of magnitudes with respect to the base β

This results in a finite number of discrete real that can be represented perfectly by the system (further significant figures are rounded off) which are denser towards the minimum order of magnitude. To prevent multiple ways of representing the same number (for example $1 \times 10 = 0.1 \times 10^2$), a **normalized** floating point system requires that $d_0 \neq 0$.

Floating point arithmetic:

1. Convert floating point numbers to the same exponent (ignoring normalized form rules)
2. Perform the operation in binary / whichever base as usual, preserving the common exponent
3. Round off any significant figures lost to precision
4. Normalize ($d_0 \neq 0$) and adjust the exponent accordingly

The total precision / exponent range is dictated by the IEEE standard for a given system, where one bit is usually reserved for signing too. **Double** uses two 32 / 64 bit words (as the name implies), allowing more precision bits and a greater exponent range than **float**; when working on a program where memory usage is non-critical, **double** is greatly preferred.

Rules of thumb:

1. Avoid equality tests involving floating point numbers - these can often return misleading results as many values (such as 1.1) are not perfectly represented in binary
2. Adding numbers of very different orders of magnitude results to lost precision